

[Get started](#)[Open in app](#)

499K Followers · About Follow

You have **2** free member-only stories left this month. [Sign up for Medium and get an extra one](#)

Comparing Modern Scalable Hyperparameter Tuning Methods

A comparison of Random search, Bayesian search using HyperOpt, Bayesian search combined with Asynchronous Hyperband and Population-Based Training



Ayush Chaurasia Sep 28 · 6 min read ★



Source — <https://pixabay.com/photos/excavators-bucket-wheel-excavators-1050501/>

In this post, we'll compare the following hyper-parameter optimization methods.

- Random Search
- Bayesian Search using HyperOpt
- Bayesian Search combined with Asynchronous Hyperband
- Population-Based Training

[Get started](#)[Open in app](#)

maximizing the inception score.

We'll use Ray Tune to perform these experiments and track the results on the W&B dashboard.

I also did a video on my channel that goes in-depth in explaining this experiment -

Comparing State of the Art Hyperparameter opt...



[Link to the Live Dashboard](#)

The Search Space

We'll use the same search space for all the experiments in order to make the comparison fair.

```
1 config = {  
2     "netG_lr": lambda: np.random.uniform(1e-2, 1e-5),  
3     "netD_lr": lambda: np.random.uniform(1e-2, 1e-5),  
4     "beta1": [0.3, 0.5, 0.8]  
5 }
```

space.py hosted with ❤ by GitHub

[view raw](#)

[Get started](#)[Open in app](#)

This will also act as the baseline metric for our comparison. Our experimental setup has 2 GPUs and 4 CPUs. We'll parallelize the operation across multiple GPUs. Ray Tune does this automatically for you if you specify the `resources_per_trial`.


```

1 analysis = tune.run(
2     dcgan_train,
3     resources_per_trial={'gpu': 1,'cpu':2}, # Tune will use this information to parallelize the
4     num_samples=10,
5     config=config
6 )

```

[randoms.py hosted with ❤ by GitHub](#) [view raw](#)

Let's see the results



[Get started](#)[Open in app](#)

As expected, we get varied results.

- **Some of the models did optimize** as the tuner got lucky and chose the right set of hyper-parameters
- but some models' **inception score graph remained flat** as they did not optimize due to bad hyper-parameter values.
- Thus, when using a random search, you might end up reaching the optimal value but **you definitely end up wasting a lot of resources** on the runs that don't add any value.

Bayesian Search With HyperOpt

The basic idea behind Bayesian Hyperparameter tuning is to **not be completely random in your choice for hyper-parameters** but instead use the **information from the prior runs** to choose the hyperparameters for the next run. Tune supports HyperOpt which implements Bayesian search algorithms. Here's how you do it.

```


1 # Step 1: Specify the search space
2 hyperopt_space= {
3     "netG_lr": hp.uniform( "netG_lr", 1e-5, 1e-2),
4     "netD_lr": hp.uniform( "netD_lr", 1e-5, 1e-2),
5     "beta1":hp.choice("beta1",[0.3,0.5,0.8])  }
6
7 #step 2: initialize the search_alg object and (optionally) set the number of concurrent runs
8 hyperopt_alg = HyperOptSearch(space = hyperopt_space,metric="is_score",mode="max")
9 hyperopt_alg = ConcurrencyLimiter(hyperopt_alg, max_concurrent=2)
10
11 #Step 3: Start the tuner
12 analysis = tune.run(
13     dcgan_train,
14     search_alg = hyperopt_alg, # Specify the search algorithm
15     resources_per_trial={'gpu': 1,'cpu':2},
16     num_samples=10,
17     config=config
18 )

```

[hyperopt.py](#) hosted with ❤ by GitHub

[view raw](#)

Here's what results look like

[Get started](#)[Open in app](#)

Inference

- There are significant improvements compared to the previous run as there is only 1 flat curve.
- This implies that the search algorithm chose the hyper-parameter values based on the results of previous runs.
- On average, the runs performed better than random search
- Resource wastage can be avoided by terminating the bad runs earlier.

Bayesian Search with Asynchronous HyperBand

The idea Asynchronous Hyperband is to **eliminate or terminate the runs that don't perform well**. It makes sense to combine this method with the Bayesian search to see

[Get started](#)[Open in app](#)

```

1  from ray.tune.schedulers import AsyncHyperBandScheduler
2  sha_scheduler = AsyncHyperBandScheduler(metric="is_score",
3                                              mode="max", max_t=300)
4
5  analysis = tune.run(
6      dcgan_train,
7      search_alg = hyperopt_alg, # Specify the search algorithm
8      scheduler = sha_scheduler, # Specify the scheduler
9      ...
10 )

```

hyperband.py hosted with ❤ by GitHub

[view raw](#)

Let us now see how this performs




Image by author

[Get started](#)[Open in app](#)

terminated earlier.

- The highest accuracy achieved was still slightly higher than the runs without the Hyperband scheduler.
- Thus, by terminating bad runs early on in the training process, we have not only speeded up the tuning job but also saved compute resources.

Population-Based training




Image source — <https://wandb.ai/wandb/DistHyperOpt/reports/Modern-Scalable-Hyperparameter-Tuning-Methods--VmIldzoyMTQxODM>

The last tuning algorithm that we'll cover is population-based training (PBT) introduced by Deepmind research. The basic idea behind the algorithm in layman terms:

- Run the optimization process for some samples for a given time step(or iterations) T.
- After every T iterations, compare the runs and copy the weights of good performing runs to the bad performing runs and change their hyper-parameter

[Get started](#)[Open in app](#)

seems simple, there is a lot of complex optimization math that goes into building this from scratch. **Tune provides a scalable and easy-to-use implementation of the SOTA PBT algorithm**

```



1 # Step 1: Initialize the PBT scheduler
2 pbt_scheduler = PopulationBasedTraining(
3     time_attr="training_iteration", # Set the time attribute as training iterations
4     metric="is_score",
5     mode="max",
6     perturbation_interval=5, #The time interval T after which you perform the PBT operation
7     hyperparam_mutations={
8         # distribution for resampling
9         "netG_lr": lambda: np.random.uniform(1e-2, 1e-5),
10        "netD_lr": lambda: np.random.uniform(1e-2, 1e-5),
11        "beta1": [0.3,0.5,0.8]
12    })
13
14 # step 2: run the tuner
15 analysis = tune.run(
16     dcgan_train,
17     scheduler=pbt_scheduler, #For using PBT
18     resources_per_trial={'gpu': 1,'cpu':2},
19     num_samples=5,
20     config=config
21 )

```

pbt.py hosted with ❤ by GitHub

[view raw](#)

Let us now look at the results.




[Get started](#)[Open in app](#)

Image by author

Inference

The results look quite surprising. There are multiple factors that are unique about these results.

- Almost all the runs have reached the optimal point
- The highest score(of 6.29) was achieved by one of the runs
- The runs that started off as bad performers or outliers have also converged as the experiment proceeded.
- There are no runs that have a flat inception score graph
- Some bad performing runs were stopped in the middle of the process
- Thus, **no resource has been wasted** on bad runs

how did PBT optimize the runs that started off with bad Hyper-parameter Values?

The answer is the **hyper-parameter mutation** done by the PBT scheduler. After every T time steps, the algorithm also mutates the values of hyper-parameters to maximize the desired metric. Here's how the parameters were mutated by the PBT scheduler for this experiment.

Hyper-parameter Mutations

Let us now see how the hyper-parameters were adjusted by the PBT algorithm to maximize the inception score


[Get started](#)[Open in app](#)

Image by author




Image by author

Inference

- The hyper-parameter values are constantly adjusted throughout the experiment.
- The runs that start off with a bad hyper-parameter value are soon updated.
- PBT operates on **explore and exploit** methodology exploring the space for good parameter values and exploits by updating the bad runs.

Reducing the Number of Runs

We'll now **reduce the number of runs** to 5 in order to make things difficult for PBT. Let's see how it performs under this restricted circumstance.

is_score

[Get started](#)[Open in app](#)

Image by author


Inference

- Even after reducing the number of runs to 5, PBT optimizer still outperforms random search and bayesian optimization.
- As expected, the bad performing runs were terminated early on in the training process

Comparing average inception scores across runs

Here's how the final comparison of the average inception scores looks like. We've averaged across 5 run sets:

- Random Search — 10 Runs (Job Type — mnist-random)
- Bayesian Search — 10 Runs (Job Type — mnist-hyperopt)
- Bayesian Search with Hyperband — 20 Runs (Job Type — mnist-SHA2-hyperopt)
- PBT scheduler — 10 Runs (Job Type — mnist-pbt2)
- PBT scheduler — 5 Runs (Job Type — mnist-pbt3)



[Get started](#)[Open in app](#)

Ending Note

If you enjoyed reading this, you can follow me on [twitter](#) to get more updates. I also make deep learning videos on my [youtube channel](#).

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

Your email

[Get this newsletter](#)

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

Machine Learning

Deep Learning

Programming

Computer Science

Math

[About](#) [Help](#) [Legal](#)

Get the Medium app

